

Neural networks

Ens'IA

Ensimag 2022-2023


November 28, 2022

The program:

- Introduction
- My first neuron
- **My first neural network**
- My first convolution
- AI challenge
- Reinforcement Learning?
- GANs?

And after? → you tell us!

Reminder




Images classification

→ We need a function f
such that $f(x) = y$


How to approximate this function? → Neural networks

Neural network



→ Succession of neuron layers

Sigmoid neuron



$$a_1, \dots, a_n \in [0, 1]$$

$$s = \sigma\left(\sum_{i=0}^n a_i * w_i + b\right) \text{ with } \sigma(x) = \frac{1}{1+e^{-x}}$$

Training loop

How to train the neuron?

For each epoch

 For each x

 Forward pass

 Loss calculation

 Backward pass

 (Accuracy calculation)

Image classification




Images of numbers between 0 and 9: 10 classes


Black & white images

28x28 size


Neural network



Neural network




Neural network



Input: 1D vector

Neural network



Input: 1D vector

Output: 1 hot encoding

Training loop

How to train the neuron?

For each epoch

 For each x

 Forward pass

 Loss calculation

 Backward pass

 (Accuracy calculation)

What about everything between the input and the output?

What about everything between the input and the output?

→ We need to find the best hidden layer/neuron count

What about everything between the input and the output?

→ We need to find the best hidden layer/neuron count

Which learning rate, batch size, epoch count and loss?

What about everything between the input and the output?

→ We need to find the best hidden layer/neuron count

Which learning rate, batch size, epoch count and loss?

→ We need to find them :3

To go further...

Back to the *sigmoid* activation function:

- Expensive to calculate
- *Vanishing Gradient*
- And more...


To go further...

Back to the *sigmoid* activation function:

- Expensive to calculate
- *Vanishing Gradient*
- And more...

In practice:

→ ReLU (*Rectified Linear Unit*)



To go further...

Back to the *backpropagation* and the gradient algorithm:

- Addition of a moment
- Adaptative η learning rate


To go further...

Back to the *backpropagation* and the gradient algorithm:

- Addition of a moment
- Adaptative η learning rate


In practice:

→ Adam



To go further...

A major problem: *overfitting*



→ *Dropout*

It works great!

We have to search for the best model

Other neural network types exist (CNN, RNN, LSTM...)

Join us on Discord !

Useful to ask questions, contact us or to pass on information! →
<https://discord.gg/UgTRbRFqNv>



To go further...

- Kaggle
- CS231N
- <http://neuralnetworksanddeeplearning.com/>
- <http://www.deeplearningbook.org/>